

int15proj_FINAL_VERSION

July 8, 2019

1 INT15 FINAL PROJECT

1.1 Due Date: June 12, midnight

Factors in Determination of Player Salaries and Win Percentages

Names: Matthew Coleman and Stephen Lantin

1.1.1 We are going to be analyzing the baseball dataset to see if we can determine batting factors which determine players salaries, as well as look at team wins and whether the ratio of runs to earned runs can be used to predict win probabilities of a team

Below are the packages used in our analysis

```
In [1]: import pandas as pd
import numpy as np

%matplotlib inline
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
from sklearn.linear_model import RidgeCV, LassoCV, Ridge, Lasso

In [2]: from sklearn.datasets import load_diabetes
from sklearn import linear_model
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error as mse

from scipy import stats
import statsmodels.api as sm
```

Reading of CSV's into dataFrames which we are going to explore and then use

```
In [3]: batting = pd.read_csv('Batting.csv')
#batting[batting['yearID'] == 2018]
fielding = pd.read_csv('FieldingOF.csv')
playerinfo = pd.read_csv('People.csv')
hof_df = pd.read_csv('HallOfFame.csv')
```

```
In [4]: batting = pd.read_csv('Batting.csv')
        #batting[batting['yearID'] == 2018]
        fieldingof = pd.read_csv('FieldingOF.csv')
        fieldingif = pd.read_csv('FieldingOF.csv')
        playerinfo = pd.read_csv('People.csv')
        #playerinfo.head()
        pitching = pd.read_csv('Pitching.csv')
        #pitching.tail(100)
        salaries = pd.read_csv('Salaries.csv')
        #salaries.head(100)
        salary = salaries['salary']
        #batting.head()
        #batting['salary'] = salary
        teams = pd.read_csv('Teams.csv')
```

Here we merge the salaries from the salary table to the batting table on player ID and year ID so we can analyze the different statistics against salary

```
In [5]: final_salaries = batting.merge(salaries, how = 'left', on = ['playerID', 'yearID'], suffixes=('_batting', '_salaries'))

        final_salaries.drop(list(final_salaries.filter(regex='_y$')), axis=1, inplace=True)
        #final_salaries = final_salaries[final_salaries['salary_x'].notnull()]
        #final_salaries['yearID'].unique()
        final_salaries = final_salaries[final_salaries['salary'].notnull()]
```

Here, we choose the salaries that are in the year 1990. We chose to do this because there was too much over plotting if we chose more than one year. We proceeded to fill all the NaN values with 0. We chose this under the assumption that if there was a blank spot, it was because they did not have any stats under there.

```
In [6]: #final_salaries
        #salaries_95 = final_salaries[(final_salaries['yearID'] >= 1990) & (final_salaries['yearID'] < 1995)]
        #salaries_00 = final_salaries[(final_salaries['yearID'] >= 1995) & (final_salaries['yearID'] < 2000)]
        # & (final_salaries['yearID'] < 1995)]
        #salaries_40 = final_salaries[(final_salaries['yearID'] >= 1940) & (final_salaries['yearID'] < 1950)]

        salaries_90 = final_salaries[(final_salaries['yearID'] == 1990)]
        salaries_90 = salaries_90.fillna(0)

        #too much over plotting, So im just going to look at the data from the 1990s. Maybe we can look at
        # 90-95, or something along those lines. Christian said its going to be hard to get any data for
        #the years as a result of overplotting.

        #NOTE:If we decide that we need to look at data across multiple years, we should look at the data for
        #want to see how other players/teams perform in relationship to other teams, we should look at the data for
        #that we can see who is the furthest team/player from the average
```

```
In [7]: #More Exploratory Code, did not use any of this in the final analysis
```

```

#standardized_salary = (salaries_95['salary'] -
#
#                        (np.sum(salaries_95['salary'])/len(salaries_95))) #/np.sqrt(len(salaries_95))

#standardized_salary
#here I am going to standardize the salary column
#(salaries_95['salary'] -np.sum(salaries_95['salary'])/len(salaries_95))

#normalize variables that depend on the number of players/team (note made by christian)

#this code below is irrelevant, I want to keep it as reference but I do not plan on using it
#std_values = (salaries_95['salary'] - (np.sum(salaries_95['salary'])/len(salaries_95)))

#salaries_95['standardized'] = std_values
#salaries_95

```

The code below was to create a heatmap based on the correlation between variables. Because we were planning to use linear regression, this was a useful diagnostic plot for determining which variables to explore


In [8]: *#Here, we look at the correlations between all the variables that will be relevant to the model. If there is anything that catches our eye, we will delve into possible relationships*

```

bat_corr = salaries_90.corr()
fig, ax = plt.subplots(figsize = (15,15))
sns.heatmap(bat_corr, annot = True, cmap = 'coolwarm')

```

Out[8]: <matplotlib.axes._subplots.AxesSubplot at 0x7ff846a019b0>

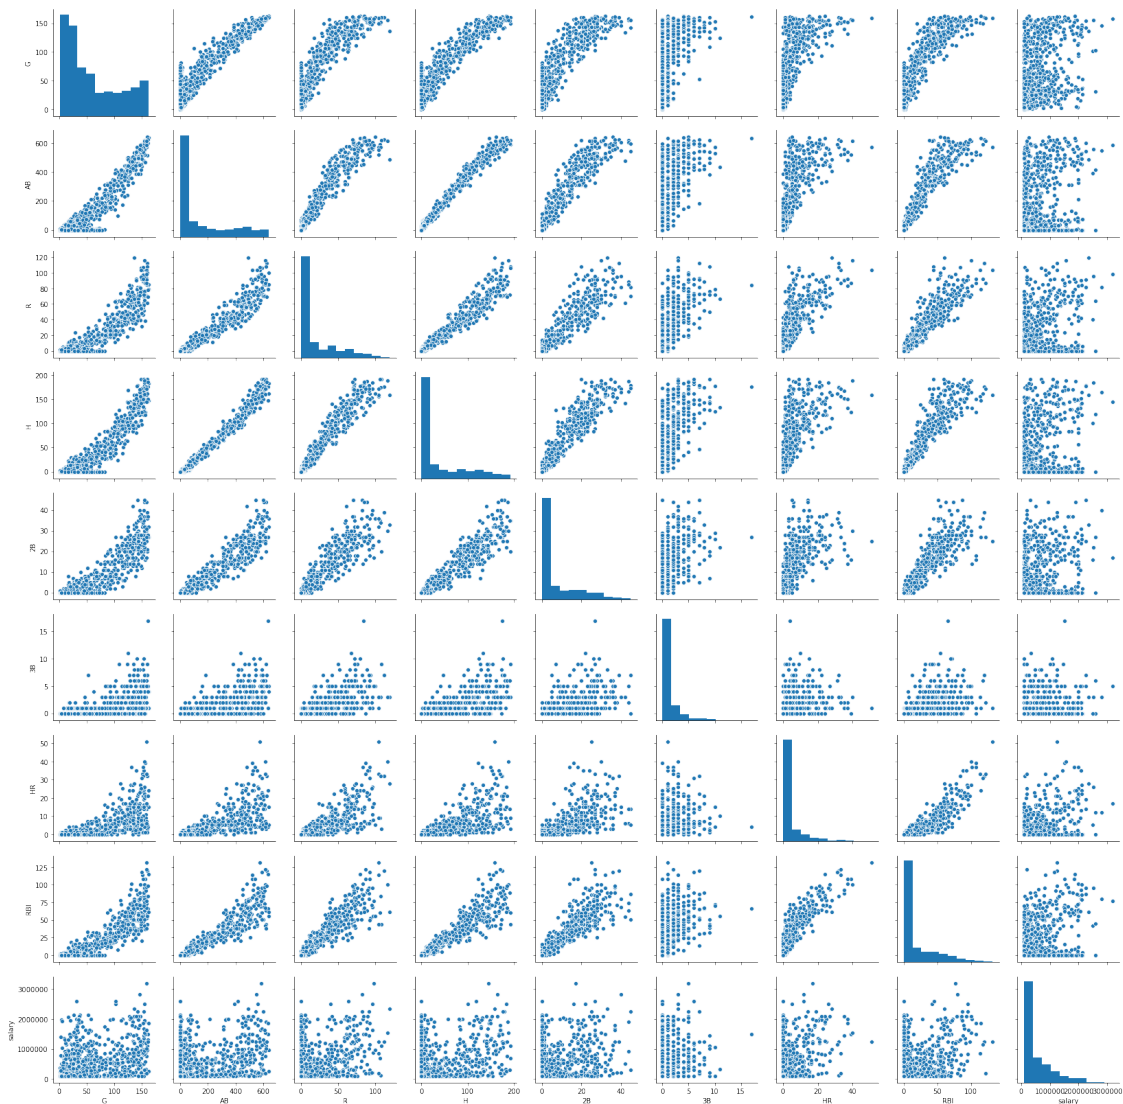


```
In [9]: #desired_columns = ['G', 'AB', 'R', 'H', '2B', '3B', 'HR', 'RBI', 'SB', 'SB', 'CS', 'BB', 'SO', 'IBB', 'HBP', 'SH', 'SF', 'GDP', 'salary']

hitting_stats = ['G', 'AB', 'R', 'H', '2B', '3B', 'HR', 'RBI', 'salary'] #, 'standardized'
other_stats = ['SB', 'SB', 'CS', 'BB', 'SO', 'salary', 'standardized']
batting_info = salaries_90.loc[:, hitting_stats]
```

Pairplot of the hitting data, visible issues of overplotting in some aspects. In stat AB through RBI compared against salary, there are clear issues with overplotting, and a line of 0 values for all of the stats. This is where we discovered that the line of zeroes was most likely pitchers, and decided to remove these data points

```
In [10]: sns.pairplot(salaries_90.loc[:, hitting_stats])
plt.show()
```



As mentioned above, looking at our analyses of the data, we noticed a “wall” of 0 values for almost variables plotted against salary. It was determined that there are many players who are not good at batting but still get paid more because of other measures of value. One of these players are pitcher, so we removed them from our analyses to look at players who could be evaluated more on their offensive aspects than their defensive ones.

```
In [11]: nonpitch_90 = salaries_90[~salaries_90['playerID'].isin(pitching['playerID'])]
         #len(sal_90_nonpitch)
         #the point of this cell is to take out all the players who are in the pitching datafr
         #variables here is pitchers salary, where they may not have as many batting statistic
         #highly because their pitching is valuable
         test_hit = nonpitch_90
         sample_50 = nonpitch_90.sort_values(by='HR',ascending = False).iloc[0:50,:]
         sample_50.shape #[['playerID','HR']]
```

```
Out[11]: (50, 23)
```

1.1.2 Here, we are splitting our data into `data_tr`, which will be our training data, and then `data_te`, which will be our testing data. By splitting the data this way, we will be able to train the model and then use the test data to test the model.

```
In [12]: data_tr, data_te = train_test_split(nonpitch_90, train_size = .75, test_size = .25)
```

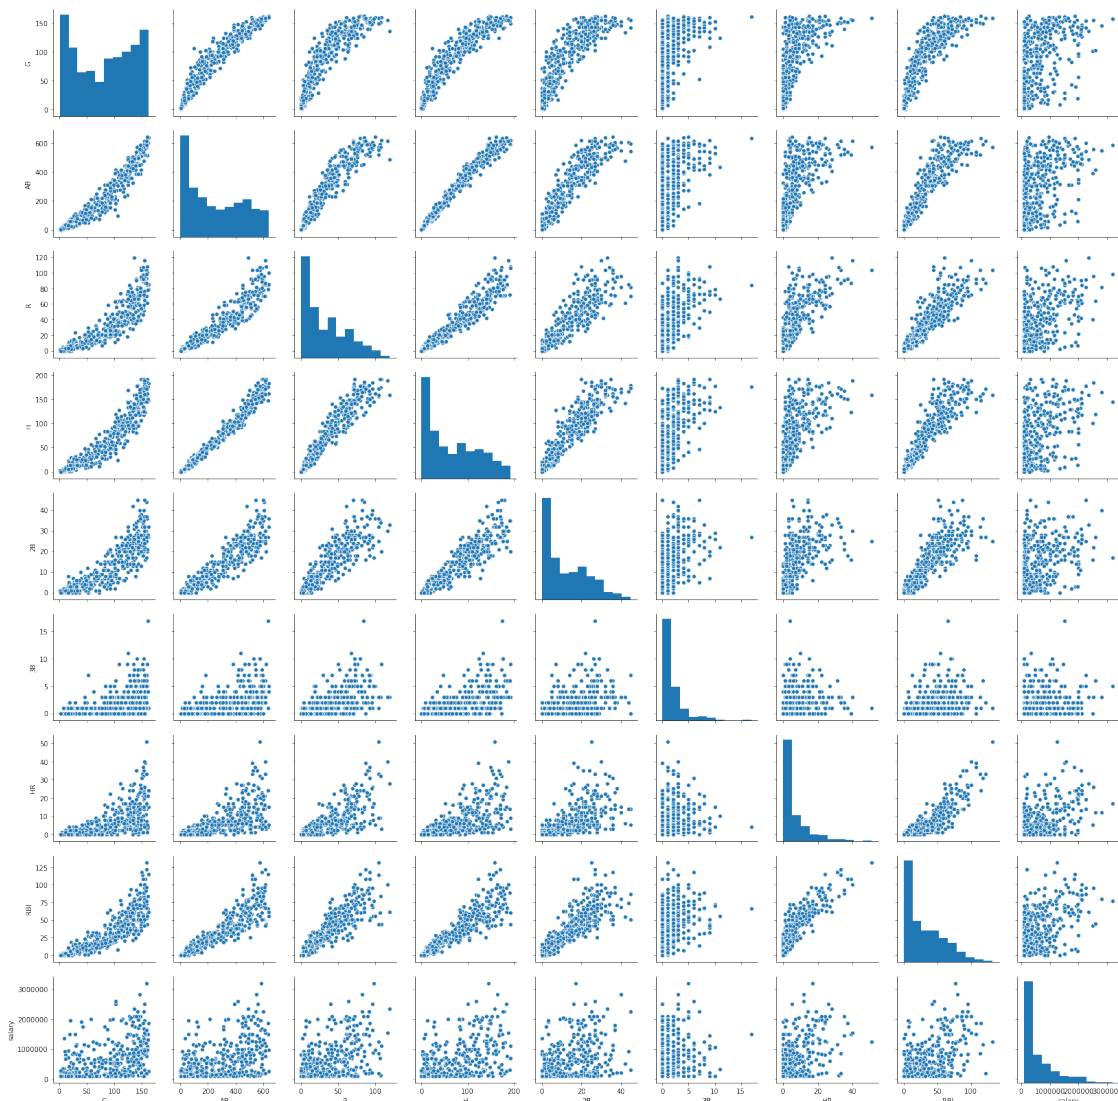
Creating a pairplot of the batting data without pitchers included shows that there is no longer an issue with the vertical line of zero values. Many of the plots have correlation, even if very weak.

```
In [13]: # We can use this pairplot to look at the relationships between variables
```

```
sns.pairplot(nonpitch_90.loc[:, hitting_stats])
```

```
plt.show()
```


```
#this pairplot is not perfect, but it is by far an improvement from the one before
```



Correlation matrix of the nonpitch data. Many of the correlations are better.

```
In [14]: nonpitchbat_corr = nonpitch_90.corr()
fig, ax = plt.subplots(figsize = (15,15))
sns.heatmap(nonpitchbat_corr, annot = True, cmap = 'coolwarm')
```

```
Out[14]: <matplotlib.axes._subplots.AxesSubplot at 0x7ff841612240>
```



We fit all the models that we are interested in below.

```
In [15]: fullmodel = sm.OLS(data_tr['salary'], data_tr[['AB', 'R', 'H', '2B', '3B', 'HR', 'RBI', 'SB',
                                                         'CS', 'BB', 'SO', 'IBB', 'HBP', 'SH']])
```



```

hit_model = sm.OLS(data_tr['salary'], data_tr[['2B', '3B', 'HR']]).fit()

nonhit = sm.OLS(data_tr['salary'], data_tr[['R', 'H', 'RBI', 'SB', 'CS', 'BB', 'IBB']]).fit()

ibb_model = sm.OLS(data_tr['salary'], data_tr['IBB']).fit()

rbi_model = sm.OLS(data_tr['salary'], data_tr['RBI']).fit()

newfull = sm.OLS(data_tr['salary'], data_tr[['AB', '2B', 'SB',
                                             'CS', 'BB', 'IBB', 'SH']]).fit()

#fullmodel.summary()
#hit_model.summary()

In [16]: fullmod_predictions = fullmodel.predict(data_te[['AB', 'R', 'H', '2B', '3B', 'HR', 'RBI', 'SB',
                                                         'CS', 'BB', 'SO', 'IBB', 'HBP', 'SH']])

hit_predictions = hit_model.predict(data_te[['2B', '3B', 'HR']])
nonhit_predictions = nonhit.predict(data_te[['R', 'H', 'RBI', 'SB', 'CS', 'BB', 'IBB']])
ibb_predictions = ibb_model.predict(data_te['IBB'])
rbi_predictions = rbi_model.predict(data_te['RBI'])
test_salaries = data_te['salary']
#fig, ax = plt.subplots(figsize = (10,10))
#sns.scatterplot(x=predictions.index, y=predictions, color = 'r')
#sns.scatterplot(x=predictions.index, y=fullmod_predictions)

```

Creating a user defined function below for giving the percentage of underestimates and over estimates

```

In [17]: ##defined function to find the percent above and percent below statistics

def overlower_estims(predictions,actual):
    difference = predictions - actual
    predict_error = []

    for value in difference:
        if value >= 0:
            predict_error.append(1)
        elif value < 0:
            predict_error.append(0)

    percent_below = (len(predict_error)-np.sum(predict_error))/len(predict_error)
    print("Percent of Under Estimates: ", 100*percent_below, ", Percent of Over Estimates: ", 100*percent_above)

```

Creating a user defined function to test the accuracy of our model at a \$100,000 tolerance level.


```
In [18]: def percent_correct(predictions, actual):
        correct_ibb = np.isclose(predictions, actual, atol = 100000)
        print("Percent of correct predictions: ", 100*(np.sum(correct_ibb)/len(correct_ibb)))
```

Percent of under estimates and correct predictions for the full model

```
In [19]: #FULL MODEL PREDICTIONS
        overlower_estims(fullmod_predictions, test_salaries)
        percent_correct(fullmod_predictions, test_salaries)
        rmsef= np.sqrt(mse(test_salaries, fullmod_predictions))
        rmsef
```

Percent of Under Estimates: 68.10344827586206 , Percent of Over Estimates: 31.89655172413793
Percent of correct predictions: 36.206896551724135

Out[19]: 389436.0183545164

Percent of under estimates and correct predictions for the hit model

```
In [20]: #HIT PREDICTIONS
        overlower_estims(hit_predictions, test_salaries)
        percent_correct(hit_predictions, test_salaries)
        rmsehit= np.sqrt(mse(test_salaries, hit_predictions))
        rmsehit
```

Percent of Under Estimates: 72.41379310344827 , Percent of Over Estimates: 27.58620689655172
Percent of correct predictions: 37.06896551724138

Out[20]: 481656.7625021634

Percent of under estimates and correct predictions for the non full model

```
In [21]: #THIS IS THE PERCENT BELOW FOR THE NON-FULL MODEL WITH NON HITTING STATS SUCH AS SB,
        overlower_estims(nonhit_predictions, test_salaries,)
        percent_correct(nonhit_predictions, test_salaries)
        rmsenf = np.sqrt(mse(test_salaries, nonhit_predictions))
        rmsenf
```

Percent of Under Estimates: 72.41379310344827 , Percent of Over Estimates: 27.58620689655172
Percent of correct predictions: 29.310344827586203

Out[21]: 410159.2936268346

Percent of under estimates and correct predictions for the intentional ball model

In [22]: *#INTENTIONAL PREDICTIONS*

```
overlower_estims(ibb_predictions, test_salaries)
percent_correct(ibb_predictions, test_salaries)
rmseibb = np.sqrt(mse(test_salaries, ibb_predictions))
rmseibb
```

Percent of Under Estimates: 80.17241379310344 , Percent of Over Estimates: 19.82758620689655
Percent of correct predictions: 31.896551724137932

Out [22]: 495166.50024517084

Percent of under estimates and correct predictions for the runs batted in model

In [23]: `overlower_estims(rbi_predictions, test_salaries)`
`percent_correct(rbi_predictions, test_salaries)`

Percent of Under Estimates: 70.6896551724138 , Percent of Over Estimates: 29.31034482758621
Percent of correct predictions: 36.206896551724135

1.2 Note: We attempted running the same analyses that we did on the batting data, except with the pitching data. We encountered many of the same problems, and overall could not find anything interesting related to what we set out to explore. As a result, we did not include any analysis relating to pitching in our report. Here is the code, showing there was exploration into pitching.

In [24]: `pitching_salaries = pitching.merge(salaries, how = 'left', on = ['playerID', 'yearID'])`
`pitching_salaries.drop(list(pitching_salaries.filter(regex='_y$')), axis=1, inplace=True)`

In [25]: `pitching_90 = pitching_salaries[pitching_salaries['yearID'] == 1990]`
`pitching_90 = pitching_90[pitching_90['salary'].notnull()]`


```
pitching_90.head(10)
pitching_90.columns
```

Out [25]: Index(['playerID', 'yearID', 'stint', 'teamID', 'lgID', 'W', 'L', 'G', 'GS',
'CG', 'SHO', 'SV', 'IPouts', 'H', 'ER', 'HR', 'BB', 'SO', 'BAOpp',
'ERA', 'IBB', 'WP', 'HBP', 'BK', 'BFP', 'GF', 'R', 'SH', 'SF', 'GIDP',
'salary'],
dtype='object')

In [26]: `pitching_90 = pitching_salaries[pitching_salaries['yearID'] == 1990]`
`pitch_tr, pitch_te = train_test_split(pitching_90, train_size = .75, test_size = .25)`
#after doing some research, it seems that the 50 ERA point is highly leveraged point
#as I could not find anything relating to a player 'Dave Martinez having a 50 ERA'
#remove the outlier

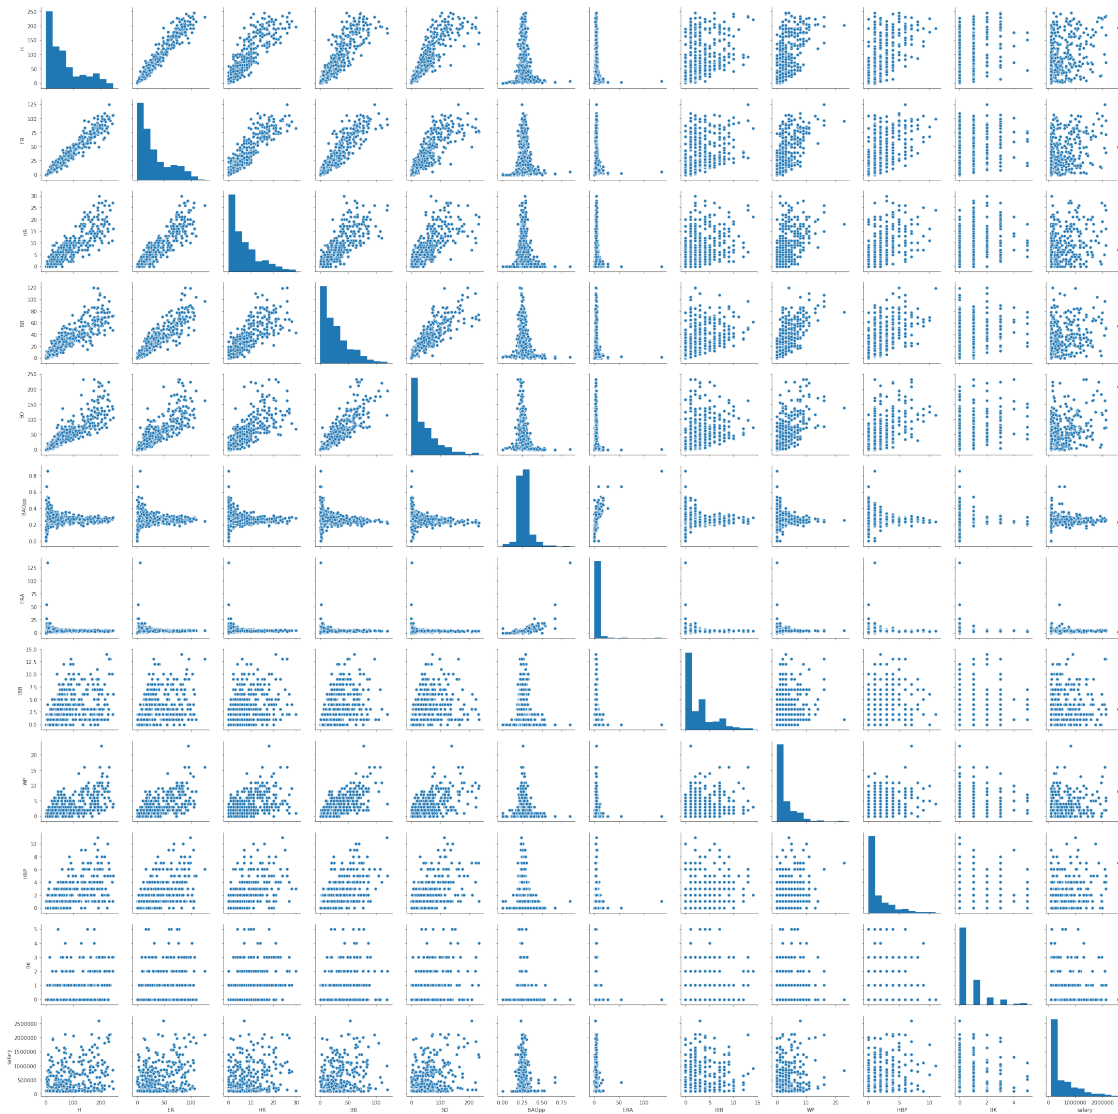
```
In [27]: fig, ax = plt.subplots(figsize = (10,10))
        sns.scatterplot(x=pitching_90['ERA'] , y=pitching_90['salary'],
                        sizes=(1, 8), linewidth=0,
                        ax=ax)
```

Out[27]: <matplotlib.axes._subplots.AxesSubplot at 0x7ff8477bcef0>




```
In [28]: sns.pairplot(pitching_90.loc[:,['H', 'ER', 'HR', 'BB', 'SO', 'BAOpp',
        'ERA', 'IBB', 'WP', 'HBP', 'BK', 'salary']])
        plt.show()
```

```
/opt/conda/lib/python3.6/site-packages/numpy/lib/histograms.py:824: RuntimeWarning: invalid va
    keep = (tmp_a >= first_edge)
/opt/conda/lib/python3.6/site-packages/numpy/lib/histograms.py:825: RuntimeWarning: invalid va
    keep &= (tmp_a <= last_edge)
```



```
In [29]: pit_corr = pitching_90[ ['H', 'ER', 'HR', 'BB', 'SO', 'BAOpp',
                                'ERA', 'IBB', 'WP', 'HBP', 'BK', 'salary']].corr()
fig, ax = plt.subplots(figsize = (20,20))
sns.heatmap(pit_corr, annot = True)
```

```
Out[29]: <matplotlib.axes._subplots.AxesSubplot at 0x7ff835769748>
```



1.2.1 End of pitching analysis

1.3 Working with team data

Choosing the team data from 1990 to 2000 and adding a fracwin column, this column is calculated as $\frac{Wins}{Wins+Losses}$

```
In [30]: teams_full = teams[(teams['yearID']>=1990) & (teams['yearID']<2000)]
         teams_90 = teams_full.loc[:, 'R':'FP']
         teams_90['W'] = teams['W']
         teams_90['L'] = teams['L']
```

```
teams_90['fracwin'] = (teams['W']/(teams['W'] + teams['L'])) # (wins/(wins + losses))
#teams_90['fracwin']
```

In [31]: *#function to add a 1 if more wins than losses.*

```
def wins(row):
    if row['fracwin'] >= .5:
        val = 1
    elif row['fracwin'] < .5:
        val = 0
    return val
```

In [32]: *#adding a row of dummy variables to whether they won most of their games or not*

```
teams_90['target'] = teams_90.apply(wins, axis = 1)
teams_90.head()
```

Out [32]:

	R	AB	H	2B	3B	HR	BB	SO	SB	CS	...	HRA	\
2047	682	5504	1376	263	26	162	473.0	1010.0	92.0	55.0	...	128	
2048	669	5410	1328	234	22	132	660.0	962.0	94.0	52.0	...	161	
2049	699	5516	1502	298	31	106	598.0	795.0	53.0	52.0	...	92	
2050	690	5570	1448	237	27	147	566.0	1000.0	69.0	43.0	...	106	
2051	682	5402	1393	251	44	106	478.0	903.0	140.0	90.0	...	106	

	BBA	SOA	E	DP	FP	W	L	fracwin	target
2047	579	938	158	133	0.974	65	97	0.401235	0
2048	537	776	93	151	0.985	76	85	0.472050	0
2049	519	997	123	154	0.980	88	74	0.543210	1
2050	544	944	142	186	0.978	80	82	0.493827	0
2051	548	914	124	169	0.980	94	68	0.580247	1

[5 rows x 30 columns]

In [33]: `win_tr, win_te = train_test_split(teams_90, train_size = .75, test_size = .25)`

`win_tr.head()`

Out [33]:

	R	AB	H	2B	3B	HR	BB	SO	SB	CS	...	HRA	\
2237	829	5628	1531	279	25	161	617.0	953.0	126.0	72.0	...	202	
2238	791	5528	1490	268	37	174	597.0	1160.0	108.0	58.0	...	111	
2243	651	5484	1386	269	27	142	518.0	1113.0	190.0	67.0	...	173	
2068	640	5474	1419	251	26	107	596.0	749.0	105.0	51.0	...	120	
2186	693	4963	1315	267	39	158	440.0	953.0	105.0	37.0	...	162	

	BBA	SOA	E	DP	FP	W	L	fracwin	target
2237	605	1050	123	140	0.980	84	78	0.518519	1
2238	450	1196	114	136	0.982	101	61	0.623457	1
2243	558	1159	106	129	0.982	76	86	0.469136	0
2068	606	1064	130	152	0.979	77	85	0.475309	0
2186	518	926	115	115	0.979	73	71	0.506944	1


[5 rows x 30 columns]

Here, we `train_test_split` the win data.

Below is the correlation plot for the team statistics The target column is going to be the target column for our PCA analysis of wins

```
In [34]: teamcorr = teams_90.corr()
fig, ax = plt.subplots(figsize = (20,20))
sns.heatmap(teamcorr, annot = True, cmap = 'coolwarm')
```

```
Out[34]: <matplotlib.axes._subplots.AxesSubplot at 0x7ff8343d9278>
```

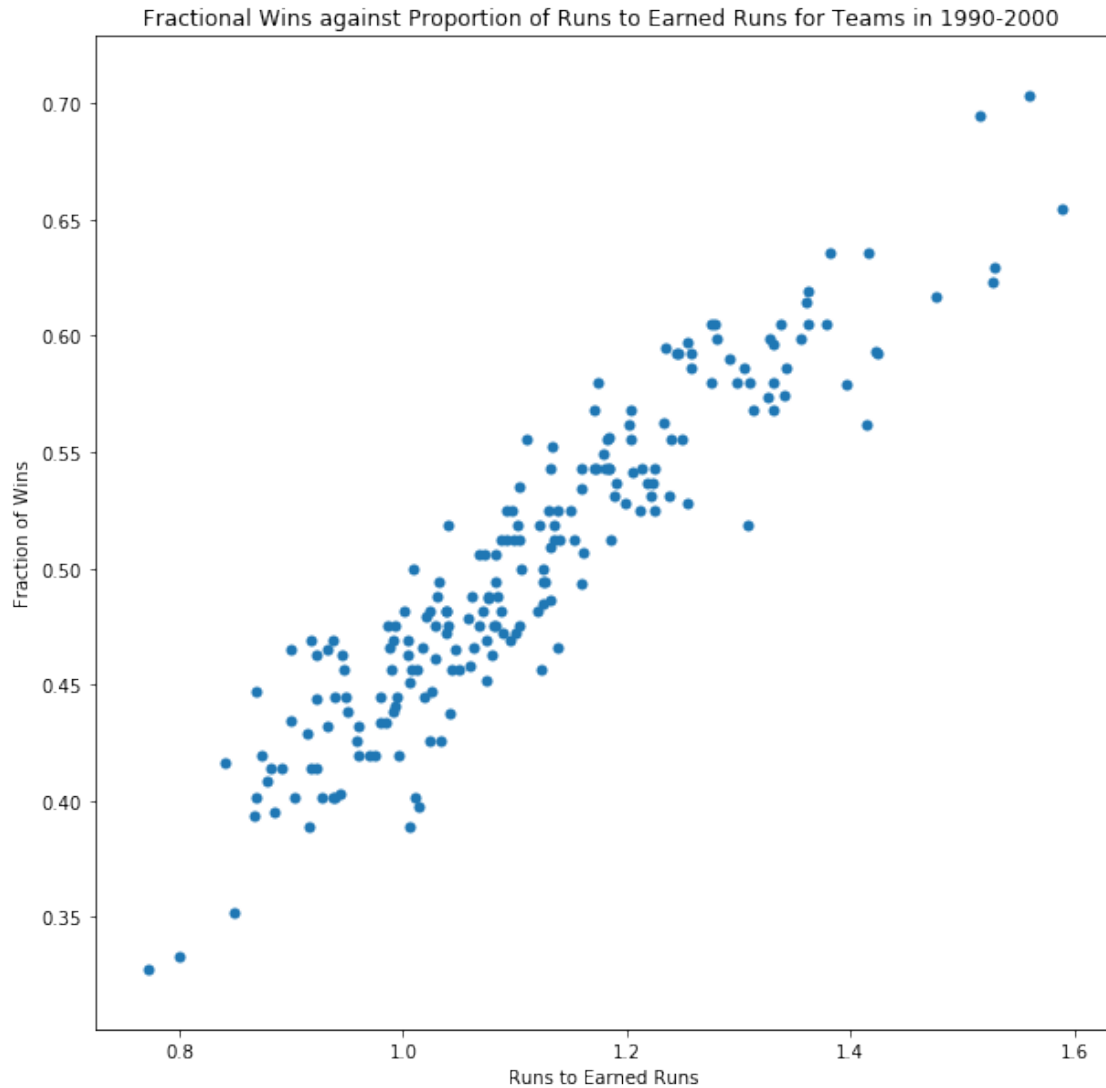


```
In [35]: fig, ax = plt.subplots(figsize = (10,10))
sns.scatterplot(x=(win_tr['R']/win_tr['ER']), y=win_tr['fracwin'],
sizes=(1, 8), linewidth=0, ax=ax)
```



```
plt.title("Fractional Wins against Proportion of Runs to Earned Runs for Teams in 1990-2000")
plt.ylabel("Fraction of Wins")
plt.xlabel("Runs to Earned Runs")
```

```
Out[35]: Text(0.5, 0, 'Runs to Earned Runs')
```



After looking at this plot, lets run a simple linear regression of this data.

```
In [36]: ratio_model = sm.OLS(win_tr['fracwin'],win_tr['R']/win_tr['ER']).fit()

ratio_model.summary()
```

```
Out[36]: <class 'statsmodels.iolib.summary.Summary'>
"""
```

OLS Regression Results

```

=====
Dep. Variable:          fracwin    R-squared:                0.997
Model:                  OLS        Adj. R-squared:           0.997
Method:                 Least Squares    F-statistic:             7.841e+04
Date:                  Tue, 18 Jun 2019    Prob (F-statistic):       5.80e-269
Time:                  18:34:59    Log-Likelihood:          463.97
No. Observations:      208    AIC:                     -925.9
Df Residuals:          207    BIC:                     -922.6
Df Model:              1
Covariance Type:       nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
x1	0.4498	0.002	280.011	0.000	0.447	0.453

```

=====
Omnibus:                 3.759    Durbin-Watson:           2.019
Prob(Omnibus):           0.153    Jarque-Bera (JB):        3.443
Skew:                    -0.308    Prob(JB):                0.179
Kurtosis:                3.133    Cond. No.                1.00
=====

```

Warnings:

```

[1] Standard Errors assume that the covariance matrix of the errors is correctly spec.
"""

```


Now, the plot of the fitted line on the scatterplot

```

In [37]: fig, ax = plt.subplots(figsize = (10,10))
sns.regplot(x=(win_tr['R']/win_tr['ER']), y=win_tr['fracwin'])
plt.title("Fractional Wins against Proportion of Runs to Earned Runs for Teams in 199")
plt.ylabel("Fraction of Wins")
plt.xlabel("Runs to Earned Runs")

Out[37]: Text(0.5, 0, 'Runs to Earned Runs')

```



```
In [38]: #Fucntion for the percent correct of the win prediction at a 5% tolerance and at a 1%
def percent_correct_wins_5(predictions, actual):
    correct_ibb = np.isclose(predictions, actual, atol = .05)
    print("Percent of correct predictions at a 5% tolerance level: ", 100*(np.sum(correct_ibb)/len(predictions)))

#Fucntion for the percent correct of the win prediction
def percent_correct_wins_1(predictions, actual):
    correct_ibb = np.isclose(predictions, actual, atol = .01)
    print("Percent of correct predictions at a 1% tolerance level: ", 100*(np.sum(correct_ibb)/len(predictions)))

In [39]: ratio_pred = ratio_model.predict((win_te['R']/win_te['ER']))
overlower_estims(ratio_pred, win_te['fracwin'])
percent_correct_wins_5(ratio_pred, win_te['fracwin'])
percent_correct_wins_1(ratio_pred, win_te['fracwin'])
```

```
rmsef= np.sqrt(mse(win_te['fracwin'], ratio_pred))
rmsef
```

```
Percent of Under Estimates:  42.857142857142854 , Percent of Over Estimates:  57.14285714285714
Percent of correct predictions at a 5% tolerance level:  87.14285714285714
Percent of correct predictions at a 1% tolerance level:  25.71428571428571
```

```
Out [39]: 0.03187252007689644
```

At a 5% tolerance, the model predicts the correct values about 90% of the time. At a 1% tolerance level, the model predicts the correct value 30 percent of the time. This means that if we want a probability for a team winning within 5% in a game based on the ratio of runs to earned runs, then our accuracy would be about 90 percent. If we wanted to predict the probability of a team winning within a margin of 1%, then we would be correct about 30 percent of the time.

Lets run some PCA on the win data First, lets separate the features from the data set

```
In [40]: features = teams_90.loc[:, 'R': 'FP'].values
        target = teams_90['target'].values
```

```
In [41]: #standardizing the features
        X = StandardScaler().fit_transform(features)
```

```
In [42]: #Performing the PCA, we will choose 2 components
        pca = PCA(n_components = 10)
```

```
principal_components = pca.fit_transform(X)
```

```
principal_comp_df = pd.DataFrame(data = principal_components, columns = ['PC 1', 'PC 2', 'PC 3', 'PC 4', 'PC 5', 'PC 6', 'PC 7', 'PC 8', 'PC 9', 'PC 10'])
```

```
principal_comp_df['target'] = target
```

```
pc_df = principal_comp_df
```

```
pc_df.head(10)
```

```
Out [42]:
```

	PC 1	PC 2	PC 3	PC 4	PC 5	PC 6	PC 7	\
0	-0.241528	0.792956	3.288464	-2.149942	1.206737	0.149560	-0.644941	
1	0.554534	0.540821	-1.259166	0.843940	0.397272	1.473680	-1.143226	
2	0.609446	-2.124669	-0.215270	-1.078822	1.804947	-0.177452	0.602859	
3	-0.083078	-1.665927	0.912958	-1.936269	2.581284	0.779326	-0.178385	
4	0.706565	-3.815984	1.185996	1.434728	-0.235867	0.319171	0.078276	
5	0.122691	-0.510835	1.586561	0.786459	0.638532	-0.519625	0.227352	
6	0.915105	-3.209060	-0.431365	0.697227	-1.109108	0.478450	0.737815	
7	-0.112757	-0.856601	0.471309	1.326847	1.062865	-1.247177	1.080311	
8	-1.153458	-0.188812	0.537086	-0.621523	1.715364	0.728353	-0.642681	

9	1.996384	-1.852817	2.759792	0.862877	-0.693166	0.557047	-1.728794
10	-0.023389	-1.606689	1.097046	0.674353	1.548812	-0.501398	1.230251
11	1.060467	-2.292746	1.038405	-0.666616	1.731053	0.256155	-1.059229
12	0.492732	-1.057358	-0.200066	1.313758	0.869539	0.348615	2.096576
13	-0.149680	-2.637619	2.367622	1.113091	2.197073	-2.033489	-0.730532
14	1.022435	-4.472712	1.043839	2.302536	-1.084662	0.793423	-1.494192
15	0.371943	0.655785	1.413534	-0.940926	0.215361	2.208592	-0.546503
16	0.923958	-2.572841	-0.929223	-2.939410	0.530267	-1.746322	-0.484294
17	1.192113	-3.568228	-3.055137	0.218459	0.108713	1.433461	-1.359412
18	0.742949	-0.053752	0.696293	-0.490131	1.946102	1.169046	-0.363708
19	0.916362	-2.689798	0.299773	0.314256	0.803609	-2.572582	-0.023840
20	0.721407	-2.128203	1.963552	-0.457587	1.373763	-0.343822	-0.196675
21	0.674815	-1.705156	0.582590	-0.841685	1.724670	-0.119424	-0.625645
22	0.298448	-0.426812	0.111003	0.974969	0.735163	0.654056	-0.019187
23	1.784843	-2.962346	2.508993	1.818220	-0.507030	-0.961532	-0.527304
24	0.196452	-1.560962	1.118155	-1.339187	1.877953	1.157298	-0.505085
25	-0.046245	-1.515784	-1.919399	2.000970	-0.219451	-0.933531	1.384335
26	0.908946	-2.788528	1.131390	-0.187408	-0.419679	-0.844568	-1.662972
27	-0.594974	1.401810	-1.439898	0.288043	1.239164	1.104782	0.846444
28	-0.402196	-0.976285	-1.115387	-0.967083	1.879576	-0.263244	0.448913
29	1.295539	-1.447216	-0.483122	-0.081563	0.348007	2.196943	0.569824
..
248	-1.044254	2.322785	-0.259047	-0.015580	0.246615	1.354741	-0.106804
249	-2.228569	-1.591418	-2.183836	-0.368687	-0.493789	-1.254281	0.997448
250	-1.282142	-1.802805	-1.691976	-0.750117	-1.791390	-0.316083	-1.117164
251	-3.247722	1.207472	-2.683074	0.820030	1.657943	1.186589	-0.098382
252	-1.809230	-1.129631	-1.277033	-1.512849	-0.848962	-2.270452	1.793160
253	-2.583877	2.028922	1.816449	0.378805	0.169361	-1.502710	0.349765
254	-2.651169	2.935160	1.658452	-1.147751	0.521254	-0.372288	0.261536
255	-2.562416	-1.183308	-1.709678	-0.057866	-2.107534	-0.010969	-0.009147
256	-5.234386	0.834023	-2.195762	0.734344	-1.057567	-1.526783	-0.949253
257	-5.597312	4.546798	0.442118	0.939531	1.837577	-0.228771	1.057820
258	-2.810135	2.832974	0.231195	0.931447	-1.869511	1.279842	0.999190
259	-2.254607	1.692646	1.896896	0.230103	-0.181941	0.023049	1.742712
260	-2.263026	-2.374923	-1.678825	0.356506	-1.146219	0.469516	-1.880767
261	-3.893247	2.643322	1.358183	1.914464	0.943381	-0.914767	2.056901
262	-2.322565	0.344118	1.028701	-0.327519	-1.220276	-0.395270	-1.914897
263	-3.775737	2.658784	0.089879	-0.719391	-0.379839	-0.894883	0.217089
264	-1.410104	1.278028	-0.308901	2.332210	0.490535	0.762113	0.751792
265	-1.608006	1.292019	3.039977	-1.761851	-1.281105	-1.187348	1.982589
266	-2.614471	-0.974045	-2.051012	-0.433407	-1.063439	-1.128687	-0.155932
267	-2.457533	-0.349108	-4.157059	1.135841	-1.303016	1.034506	-1.997872
268	-3.596756	1.968934	-1.606319	-1.888181	-0.711893	0.091011	-1.141570
269	-3.031764	1.824557	-0.677484	0.579179	-0.089962	0.224631	1.127452
270	-2.531939	0.926230	1.854178	-1.529688	-0.861982	0.046468	0.547211
271	-1.516847	0.338494	1.116268	-0.115722	-1.769162	0.679090	-2.227469
272	-4.165657	2.610551	-0.475308	-0.078993	0.318670	0.560704	-1.581306
273	-3.567161	1.867596	-1.550686	-0.142904	-0.956396	0.716564	-1.403523

274	-3.210781	1.650820	1.003853	-0.980234	-0.955154	0.302016	-1.094609
275	-3.809440	1.992402	1.391956	-0.615517	0.203356	0.714934	0.496630
276	-4.078046	0.920309	-1.180733	0.720502	0.635725	-1.970713	-0.809786
277	-4.053575	1.632746	-1.789419	-0.404834	-0.312709	0.897971	-0.347663

	PC 8	PC 9	PC 10	target
0	0.295965	0.780602	0.496861	0
1	-1.343863	-0.277328	0.105604	0
2	-1.160820	-0.830742	-0.263305	1
3	-0.777327	-0.707167	1.102157	0
4	-1.501927	-0.935459	1.169449	1
5	-0.602075	0.305665	-0.391223	0
6	-0.008512	0.468905	-0.234740	1
7	-1.049337	-0.107014	-0.524644	0
8	-0.962033	-0.096444	1.384101	0
9	-0.214086	0.720931	-0.080748	0
10	0.274101	0.086862	-0.726848	0
11	2.087207	1.237848	0.003249	1
12	-0.295866	-1.407590	-0.237952	0
13	0.951113	-0.847759	0.557475	0
14	0.262617	1.520731	0.400775	1
15	-0.203454	-0.896631	0.180811	0
16	0.806471	0.314014	-0.807455	1
17	-0.943133	-0.649688	1.397063	1
18	-0.603527	0.137909	-0.529105	0
19	0.105423	0.460856	0.297956	1
20	0.615441	0.482376	0.289259	0
21	0.152796	-1.039225	-1.044365	0
22	-1.173677	0.444479	0.754616	1
23	-0.306868	0.621982	-0.907105	0
24	0.645305	0.042559	-0.230549	1
25	-1.654882	1.080738	0.333435	1
26	0.278137	0.302532	0.625343	1
27	-1.939746	-0.132946	0.201695	0
28	-0.840905	-1.036947	-0.555893	1
29	-0.885690	-0.139131	0.220980	1
..
248	-1.219112	-0.087364	-0.551481	0
249	1.005804	1.779031	0.333884	1
250	0.704461	0.035585	-0.172356	1
251	1.397272	-0.886407	-0.094394	0
252	-0.432595	-0.304077	0.116114	1
253	-0.863784	0.233310	-0.200168	0
254	-0.043240	1.710931	0.445044	0
255	-0.439117	1.230546	0.438291	1
256	0.040387	0.490247	-0.467562	1
257	-0.111656	0.674477	0.277142	0
258	1.196173	-0.280682	0.558777	0

259	-0.132027	0.199673	-0.925812	0
260	0.553541	-0.669197	0.044001	1
261	0.795963	-0.372068	0.560166	0
262	0.896235	0.320924	-0.256060	0
263	-0.777229	-0.048127	-0.606024	0
264	0.668129	0.002507	-1.494993	0
265	-0.695438	0.224201	0.725232	0
266	-0.329988	0.362460	0.192491	1
267	-0.379441	0.073311	-1.589330	1
268	-0.484582	-1.035115	1.685102	1
269	0.451602	1.925917	-0.215019	0
270	0.327784	-0.200252	0.730195	0
271	-0.392132	0.682422	-0.144976	0
272	-0.418458	0.307544	0.635986	0
273	-0.013750	-0.307083	-0.408206	1
274	-0.213238	0.215082	0.260386	0
275	-0.925259	-1.727557	-0.169179	0
276	-0.900832	0.132112	0.677522	1
277	1.782907	-1.158332	0.231105	1

[278 rows x 11 columns]

Before we plot the principal components, let's take a look at the variance explained by the principal components


```
In [43]: variance=np.var(principal_components,axis=0)
variance_ratio = variance/np.sum(variance)
component_no = [1,2,3,4,5,6,7,8,9,10]
print(variance_ratio)
```

```
[0.417515    0.18038884 0.10079889 0.06489423 0.05699299 0.04632148
 0.04169407 0.03711056 0.0314276  0.02285634]
```

Looking at the percent variance explained by each principle component, you can see that the first component explains 69.8% of the variation. Meanwhile, the second component explains 30.1% of the variation. The scree plot below demonstrates this.

```
In [44]: sns.barplot(component_no,variance_ratio, color = 'steelblue').set(xlabel = "Principal
```


```
Out[44]: [Text(0, 0.5, 'Variance Explained'), Text(0.5, 0, 'Principal Component')]
```

Now, we can use the `pc_df` to visualize the principal components.

```
In [45]: fig = plt.figure(figsize = (8,8))
ax = fig.add_subplot(1,1,1)
ax.set_xlabel('Principal Component 1')
ax.set_ylabel('Principal Component 2')
ax.set_title('PC 1 and PC2 of Win Data', fontsize = 20)
targets = [0,1]
colors = ['r', 'g']
for target, color in zip(targets,colors):
    indicesToKeep = pc_df['target'] == target
    ax.scatter(pc_df.loc[indicesToKeep, 'PC 1']
               , pc_df.loc[indicesToKeep, 'PC 2']
               , c = color
               , s = 50)
ax.legend(targets,labels = ['More Losses', 'More Wins'])
ax.grid()
```


/opt/conda/lib/python3.6/site-packages/ipykernel_launcher.py:14: UserWarning: You have mixed p



Observing the graph of the two principal components, there is a clear distinguishing of teams with more wins and teams with more losses on the line $y = -x - 2$. On the bottom portion of the line, there seems to be teams with more wins. Above the line there are more teams that have more losses. There is another grouping of points in the upper right corner from the bigger grouping of points. The extremely interesting part about this is that the groupings of points are in the same layout at the bigger grouping. The line $y = \frac{5}{6}x + \frac{25}{3}$ divides the teams with more wins below the line, while the teams above the line are mostly teams with more losses. After some investigation, there were four expansion teams in the 1990's: Colorado Rockies(1993), Florida Marlins(1993) (Now Miami Marlins), Arizona Diamondbacks(1998), and Tampa Bay Devil Rays(1998) (Now Tampa Bay Rays). It is possible that because the analyzed data was from 1990 to 2000, this grouping represents these teams. Unfortunately, it is impossible to definitely say based on this visualization.

```
In [46]: plt.matshow(pca.components_, cmap='coolwarm')
          plt.xticks([0,1], ['1st Component', '2nd Component'], fontsize=10)
```

```
plt.colorbar()
plt.xticks(range(len(teams_90.loc[:, 'R': 'FP'].columns)), teams_90.loc[:, 'R': 'FP'].columns)
plt.show()
```



The plot above shows how much of each variable accounted for each principal component. Noticeably, the second component seems to have been based highly off of Earned Run Average (ERA), Run Average (RA), Earned Runs (ER), Home Runs (HR), and Home Runs Allowed (HRA).

The ratio of runs to earned runs seems to be a promising for a regression of Fractional wins on R/ER

1.4 Conclusions and Future Work (200 words)

In this section you should summarize your findings based on your final model in clearly understandable, non-statistical terms. What is the main message produced by your analysis? There may also be additional questions that arise, problems you encounter, or possible extensions of your analysis that could be addressed here.

Include any final comments and thoughts about your project. For example, do you trust your results? How general are your results, to what situations do they apply? Add any other comments that are relevant.

1.5 Grading Criteria

Your grade on the project will be based on the following criteria:

1. Compatibility of Scientific Question and Analysis

Is the scientific question being addressed actually of interest, and were suitable tools employed?

2. Coherent Thought Process and Presentation of Results

- Does the analysis indicate a sound understanding of methods discussed in class? This is often best judged by the preliminary comments on the questions of interest as well as the conclusions made after the analysis.

- Is the analysis presented in a clear, consistent, coherent style with the appropriately labeled requested components and visualizations?

3. Scope of the analysis and methods used

Did the analysis demonstrate a wide understanding of methodology and ideas presented throughout the quarter?

4. Reproducible results

Is it possible to reproduce the analysis and the visuals by executing the provided code?

1.6 Submission

You are required to submit two files:

1. Submit your completed writeup as a PDF to gradescope. You should address all of the components described above, adhering to the page limit, and include any figures and tables that are necessary. (Make sure to number figures and tables and include informative captions.)
2. Submit a complete jupyter notebook with all of your analyses to the okpy server. For your submission, use **this jupyter notebook** as a template (remove the instructions, replacing them with your analysis). We should be able to reproduce all of your results by running your notebook.

Before you submit the notebook, make sure that you select from the top menu Kernel -> Restart & Clear Output followed by Cell -> Run All. Verify that all computations execute correctly. There should be no errors when we run your notebook.

```
In [47]: # These lines help load your submission for grading.
         from client.api.notebook import Notebook
         ok = Notebook('final-project.ok')
         _ = ok.auth(inline=True)
```

```
=====
Assignment: final-project
OK, version v1.14.15
=====
```

Successfully logged in as matthewbcoleman@ucsb.edu

```
In [ ]: _ = ok.submit()
```

```
<IPython.core.display.Javascript object>
```

```
In [ ]:
```